

Practical R for Epidemiologists

Mark Myatt

2018-04-21

Contents

1	Introduction	5
1.1	Introducing R	5
1.2	Retrieving data	7
2	Getting acquainted with R	9
2.1	Summary	31
3	Manipulating objects and creating new functions	33

Chapter 1

Introduction

1.1 Introducing R

R is a system for data manipulation, calculation, and graphics. It provides:

- Facilities for data handling and storage
- A large collection of tools for data analysis
- Graphical facilities for data analysis and display
- A simple but powerful programming language

R is often described as an environment for working with data. This is in contrast to a *package* which is a collection of very specific tools. R is not strictly a statistics system but a system that provides many classical and modern statistical procedures as part of a broader data-analysis tool. This is an important difference between R and other statistical systems. In R a statistical analysis is usually performed as a series of steps with intermediate results being stored in objects. Systems such as SPSS and SAS provide copious output from (e.g.) a regression analysis whereas R will give minimal output and store the results of a fit for subsequent interrogation or use with other R functions. This means that R can be tailored to produce exactly the analysis and results that you want rather than produce an analysis designed to fit all situations.

R is a language based product. This means that you interact with R by typing commands such as:

```
table(SEX, LIFE)
```

rather than by using menus, dialog boxes, selection lists, and buttons. This may seem to be a drawback but it means that the system is considerably more flexible than one that relies on menus, buttons, and boxes. It also means that every stage of your data management and analysis can be recorded and edited and re-run at a later date. It also provides an audit trail for quality control purposes.

R is available under UNIX (including Linux), the Macintosh operating system OS X, and Microsoft Windows. The method used for starting R will vary from system to system. On UNIX systems you may need to issue the R command in a terminal session or click on an icon or menu option if your system has a windowing system. On Macintosh systems R will be available as an application but can also be run in a terminal session. On Microsoft Windows systems there will usually be an icon on the Start menu or the desktop.

R is an open source system and is available under the *GNU general public license* (GPL) which means that it is available for free but that there are some restrictions on how you are allowed to distribute the system and how you may charge for bespoke data analysis solutions written using the R system. Details of the general public license are available from <http://www.gnu.org/copyleft/gpl.html>.

R is available for download from <http://www.r-project.org/>.

This is also the best place to get extension packages and documentation. You may also subscribe to the R mailing lists from this site. R is supported through mailing lists. The level of support is at least as good as for commercial packages. It is typical to have queries answered in a matter of a few hours.

Even though R is a free package it is more powerful than most commercial packages. Many of the modern procedures found in commercial packages were first developed and tested using R or **S-Plus** (the commercial equivalent of R).

When you start R it will issue a prompt when it expects user input. The default prompt is:

```
>
```

This is where you type commands that call functions that instruct R to (e.g.) read a data file, recode data, produce a table, or fit a regression. For example:

```
> table(SEX, LIFE)
```

If a command you type is not complete then the prompt will change to:

```
+
```

on subsequent lines until the command is complete:

```
> table(  
+ SEX, LIFE +)
```

The > and + prompts are not shown in the example commands in the rest of this material.

The example commands in this material are often broken into shorter lines and indented for ease of understanding. The code still works as lines are split in places where R knows that a line is not complete. For example:

```
table(SEX,  
      LIFE)
```

could be entered on a single line as:

```
table(SEX, LIFE)
```

In this example R knows that the command is not complete until the brackets are closed. The following example could also be written on one line:

```
salex.lreg.coeffs <-  
  coef(summary(salex.lreg))
```

In this case R knows that the <- operator at the end of the first line needs further input.

R maintains a history of previous commands. These can be recalled and edited using the up and down arrow keys.

Output that has scrolled off the top of the output / command window can be recalled using the window or terminal scroll bars.

Output can be saved using the `sink()` function with a file name: `sink("results.out")` to start recording output. Use the `sink()` function without a file name to stop recording output: `sink()`

You can also use clipboard functions such as copy and paste to (e.g.) copy and then paste selected chunks of output into an editor or word processor running alongside R.

All the sample data files used in the exercises in this manual are space delimited text files using the general format:

```
ID AGE IQ  
1 39 94  
2 41 89
```

```

3 42 83
4 30 99
5 35 94
6 44 90
7 31 94
8 39 87

```

R has facilities for working with files in different formats including (through the use of extension packages) **ODBC** (open database connectivity) and **SQL** data sources, **EpiInfo**, **EpiData**, **Minitab**, **SPSS**, **SAS**, **S-Plus**, and **Stata** format files.

1.2 Retrieving data

All of the exercises in this manual assume that the necessary data files are located in the current working directory. All of the data files that you require to follow this material are in a ZIP archive that can be downloaded from:

<http://www.brixtonhealth.com/prfe/prfe.zip>

A command such as:

```
read.table("data/fem.dat", header = TRUE)
```

retrieves the data stored in the file named `fem.dat` which is stored in the `data` folder.

To retrieve data that is stored in files outside a different directory you need to specify the full path to the file. For example:

```
read.table("~/example/fem.dat", header = TRUE)
```

will retrieve the data stored in the file named `fem.dat` stored in the `example` directory under the user's home directory on UNIX, Linux, and OS X systems.

R follows many UNIX operating and naming conventions including the use of the backslash (`\`) character to specify special characters in strings (e.g. using `\n` to specify a new line in printed output). Windows uses the backslash (`\`) character to separate directory and file names in paths. This means that Windows users need to escape any backslashes in file paths using an additional backslash character. For example:

```
read.table("c:\\example\\fem.dat", header = TRUE)
```

will retrieve the data that is stored in the file named `fem.dat` which is stored in the `example` directory off the root directory of the `C:` drive. The Windows version of R also allows you to specify UNIX-style path names (i.e. using the forward slash (`/`) character as a separator in file paths). For example:

```
read.table("c:/example/fem.dat", header = TRUE)
```

Path names may include shortcut characters such as:

.	The current working directory
..	Up one level in the directory tree
~	The user's home directory (on UNIX-based systems)

R also allows you to retrieve files from any location that may be represented by a standard **uniform resource locator** (URL) string. For example:

```
read.table("file://~/example/fem.dat", header = TRUE)
```

will retrieve the data stored in the file named `fem.dat` stored in the `example` directory under the users home

directory on UNIX-based systems.

All of the data files used in this section are stored in the `/data` directory of this guide's GitLab repository (<https://git.validmeasures.org/datahub/datahubguide/tree/master/data>). This means, for example, that you can use the `read.table()` function specifying

`"https://git.validmeasures.org/datahub/datahubguide/tree/master/data/fem.dat"`

as the `URL` to retrieve the data that is stored in the file named `fem.dat` which is stored in the `/data` directory of this guide's GitLab repository.

Chapter 2

Getting acquainted with R

In this exercise we will use R to read a dataset and produce some descriptive statistics, produce some charts, and perform some simple statistical inference. The aim of the exercise is for you to become familiar with R and some basic R functions and objects.

The first thing we will do, after starting R, is issue a command to retrieve an example dataset:

```
fem <- read.table("fem.dat", header = TRUE)
```

This command illustrates some key things about the way R works.

We are instructing R to assign (using the `<-` operator) the output of the `read.table()` function to an object called `fem`.

The `fem` object will contain the data held in the file `fem.dat` as an R data.frame object:

```
class(fem)
```

```
## [1] "data.frame"
```

You can inspect the contents of the `fem` data.frame (or any other R object) just by typing its name:

```
fem
```

```
##   ID AGE IQ ANX DEP SLP SEX LIFE   WT
## 1  1  39 94   2   2   2   1   1  2.23
## 2  2  41 89   2   2   2   1   1  1.00
## 3  3  42 83   3   3   2   1   1  1.82
## 4  4  30 99   2   2   2   1   1 -1.18
## 5  5  35 94   2   1   1   1   2 -0.14
## 6  6  44 90  NA   1   2   2   2  0.41
```

Note that the `fem` object is built from other objects. These are the named vectors (columns) in the dataset:

```
names(fem)
```

```
## [1] "ID"    "AGE"   "IQ"    "ANX"   "DEP"   "SLP"   "SEX"   "LIFE"  "WT"
```

The `[1]` displayed before the column names refers to the numbered position of the first name in the output. These positions are known as indexes and can be used to refer to individual items. For example:

```
names(fem)[1]
```

```
## [1] "ID"
```

```
names(fem)[8]
```

```
## [1] "LIFE"
```

```
names(fem)[2:4]
```

```
## [1] "AGE" "IQ" "ANX"
```

The data consist of 118 records:

```
nrow(fem)
```

```
## [1] 118
```

each with nine variables:

```
ncol(fem)
```

```
## [1] 9
```

for female psychiatric patients.

The columns in the dataset are:

ID	Patient ID
AGE	Age in years
IQ	IQ score
ANX	Anxiety (1=none, 2=mild, 3=moderate, 4=severe)
DEP	Depression (1=none, 2=mild, 3=moderate or severe)
SLP	Sleeping normally (1=yes, 2=no)
SEX	Lost interest in sex (1=yes, 2=no)
LIFE	Considered suicide (1=yes, 2=no)
WT	Weight change (kg) in previous 6 months

The first ten records of the `fem` data.frame are:

```
##      ID AGE  IQ ANX DEP SLP SEX LIFE    WT
## 1    1  39  94   2   2   2   1   1  2.23
## 2    2  41  89   2   2   2   1   1  1.00
## 3    3  42  83   3   3   2   1   1  1.82
## 4    4  30  99   2   2   2   1   1 -1.18
## 5    5  35  94   2   1   1   1   2 -0.14
## 6    6  44  90  NA   1   2   2   2  0.41
## 7    7  31  94   2   2  NA   1   1 -0.68
## 8    8  39  87   3   2   2   1   2  1.59
## 9    9  35 -99   3   2   2   1   1 -0.55
## 10  10  33  92   2   2   2   1   1  0.36
```

You may check this by asking R to display all columns of the first ten records in the `fem` data.frame:

```
fem[1:10, ]
```

```
##      ID AGE  IQ ANX DEP SLP SEX LIFE    WT
## 1    1  39  94   2   2   2   1   1  2.23
## 2    2  41  89   2   2   2   1   1  1.00
## 3    3  42  83   3   3   2   1   1  1.82
## 4    4  30  99   2   2   2   1   1 -1.18
## 5    5  35  94   2   1   1   1   2 -0.14
## 6    6  44  90  NA   1   2   2   2  0.41
```

```
## 7 7 31 94 2 2 NA 1 1 -0.68
## 8 8 39 87 3 2 2 1 2 1.59
## 9 9 35 -99 3 2 2 1 1 -0.55
## 10 10 33 92 2 2 2 1 1 0.36
```

The space after the comma is optional. You can think of it as a *placeholder* for where you would specify the indexes for columns you wanted to display. For example:

```
fem[1:10,2:4]
```

displays the first ten rows and the second, third and fourth columns of the `fem` data.frame:

```
##    AGE  IQ ANX
## 1   39  94  2
## 2   41  89  2
## 3   42  83  3
## 4   30  99  2
## 5   35  94  2
## 6   44  90 NA
## 7   31  94  2
## 8   39  87  3
## 9   35 -99  3
## 10  33  92  2
```

NA is a special value meaning *not available* or *missing*.

You can access the contents of a single column by name:

```
fem$IQ
```

```
## [1] 94 89 83 99 94 90 94 87 -99 92 92 94 91 86 90 -99 91
## [18] 82 86 88 97 96 95 87 103 -99 91 87 91 89 92 84 94 92
## [35] 96 96 86 92 102 82 92 90 92 88 98 93 90 91 -99 92 92
## [52] 91 91 86 95 91 96 100 99 89 89 98 98 103 91 91 94 91
## [69] 85 92 96 90 87 95 95 87 95 88 94 -99 -99 87 92 86 93
## [86] 92 106 93 95 95 92 98 92 88 85 92 84 92 91 86 92 89
## [103] -99 96 97 92 92 98 91 91 89 94 90 96 87 86 89 -99
```

```
fem$IQ[1:10]
```

```
## [1] 94 89 83 99 94 90 94 87 -99 92
```

The `$` sign is used to separate the name of the data.frame and the name of the column of interest. Note that R is case-sensitive so that `IQ` and `iq` are *not* the same.

You can also access rows, columns, and individual cells by specifying row and column positions. For example, the `IQ` column is the third column in the `fem` data.frame:

```
fem[,3]
```

```
## [1] 94 89 83 99 94 90 94 87 -99 92 92 94 91 86 90 -99 91
## [18] 82 86 88 97 96 95 87 103 -99 91 87 91 89 92 84 94 92
## [35] 96 96 86 92 102 82 92 90 92 88 98 93 90 91 -99 92 92
## [52] 91 91 86 95 91 96 100 99 89 89 98 98 103 91 91 94 91
## [69] 85 92 96 90 87 95 95 87 95 88 94 -99 -99 87 92 86 93
## [86] 92 106 93 95 95 92 98 92 88 85 92 84 92 91 86 92 89
## [103] -99 96 97 92 92 98 91 91 89 94 90 96 87 86 89 -99
```

```
fem[9, ]
```

```
## ID AGE IQ ANX DEP SLP SEX LIFE WT
```

```
## 9 9 35 -99 3 2 2 1 1 -0.55
```

```
fem[9,3]
```

```
## [1] -99
```

There are missing values in the IQ column which are all coded as **-99**. Before proceeding we must set these to the special NA value:

```
fem$IQ[fem$IQ == -99] <- NA
```

The term inside the square brackets is also an index. This type of index is used to refer to subsets of data held in an object that meet a particular condition. In this case we are instructing R to set the contents of the IQ variable to NA if the contents of the IQ variable is **-99**.

Check that this has worked:

```
fem$IQ
```

```
## [1] 94 89 83 99 94 90 94 87 NA 92 92 94 91 86 90 NA 91
## [18] 82 86 88 97 96 95 87 103 NA 91 87 91 89 92 84 94 92
## [35] 96 96 86 92 102 82 92 90 92 88 98 93 90 91 NA 92 92
## [52] 91 91 86 95 91 96 100 99 89 89 98 98 103 91 91 94 91
## [69] 85 92 96 90 87 95 95 87 95 88 94 NA NA 87 92 86 93
## [86] 92 106 93 95 95 92 98 92 88 85 92 84 92 91 86 92 89
## [103] NA 96 97 92 92 98 91 91 89 94 90 96 87 86 89 NA
```

We can now compare the groups who have and have not considered suicide. For example:

```
by(fem$IQ, fem$LIFE, summary)
```

Look at the help for the `by()` function:

```
help(by)
```

Note that you may use `?by` as a shortcut for `help(by)`.

The `by()` function applies another function (in this case the `summary()` function) to a column in a data.frame (in this case `fem$IQ`) split by the value of another variable (in this case `fem$LIFE`).

It can be tedious to always have to specify a data.frame each time we want to use a particular variable. We can fix this problem by ‘attaching’ the data.frame:

```
attach(fem)
```

```
## The following objects are masked from fem (pos = 4):
```

```
##
```

```
## AGE, ANX, DEP, ID, IQ, LIFE, SEX, SLP, WT
```

We can now refer to the columns in the `fem` data.frame without having to specify the name of the data.frame. This time we will produce summary statistics for `WT` by `LIFE`:

```
by(WT, LIFE, summary)
```

```
## LIFE: 1
##   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.   NA's
## -2.2300 -0.2700  1.0000  0.7867  1.7300  3.7700     4
## -----
## LIFE: 2
##   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.   NA's
## -1.6800 -0.4500  0.6400  0.6404  1.5000  2.9500     7
```

We can view the same data as a box and whisker plot:

```
boxplot(WT ~ LIFE)
```



We can add axis labels and a title to the graph:

```
boxplot(WT ~ LIFE,
        xlab = "Life",
        ylab = "Weight",
        main = "Weight BY Life")
```

Weight BY Life




A more descriptive title might be “Weight Change BY Considered Suicide”.

The groups do not seem to differ much in their medians and the distributions appear to be reasonably


symmetrical about their medians with a similar spread of values.

We can look at the distribution as histograms:

```
hist(WT[LIFE == 1])
```




```
hist(WT[LIFE == 2])
```

Histogram of WT[LIFE == 2]

and check the assumption of normality using quantile-quantile plots:

```
qqnorm(WT[LIFE == 1])  
qqline(WT[LIFE == 1])
```

Normal Q-Q Plot

```
qqnorm(WT[LIFE == 2])
qqline(WT[LIFE == 2])
```



or by using a formal test:

```
shapiro.test(WT[LIFE == 1])
```

```
##
##  Shapiro-Wilk normality test
##
## data:  WT[LIFE == 1]
## W = 0.98038, p-value = 0.4336
```

```
shapiro.test(WT[LIFE == 2])
```

```
##
##  Shapiro-Wilk normality test
##
## data:  WT[LIFE == 2]
## W = 0.97155, p-value = 0.3292
```

Remember that we can use the `by()` function to apply a function to a data.frame, including statistical functions such as `shapiro.test()`:

```
by(WT, LIFE, shapiro.test)
```

```
## LIFE: 1
##
##  Shapiro-Wilk normality test
##
## data:  dd[x, ]
## W = 0.98038, p-value = 0.4336
```



```
##
## -----
## LIFE: 2
##
## Shapiro-Wilk normality test
##
## data:  dd[, ]
## W = 0.97155, p-value = 0.3292
```

We can also test whether the variances differ significantly using *Bartlett's test* for the homogeneity of variances:

```
bartlett.test(WT, LIFE)
```

```
##
## Bartlett test of homogeneity of variances
##
## data:  WT and LIFE
## Bartlett's K-squared = 0.32408, df = 1, p-value = 0.5692
```

There is no significant difference between the two variances.

Many functions in R have a *formula interface* that may be used to specify multiple variables and the relations between multiple variables. We could have used the formula interface with the `bartlett.test()` function:

```
bartlett.test(WT ~ LIFE)
```

```
##
## Bartlett test of homogeneity of variances
##
## data:  WT by LIFE
## Bartlett's K-squared = 0.32408, df = 1, p-value = 0.5692
```

Having checked the normality and homogeneity of variance assumptions we can proceed to carry out a t-test:

```
t.test(WT ~ LIFE, var.equal = TRUE)
```

```
##
## Two Sample t-test
##
## data:  WT by LIFE
## t = 0.59869, df = 104, p-value = 0.5507
## alternative hypothesis: true difference in means is not equal to 0
## 95 percent confidence interval:
##  -0.3382365  0.6307902
## sample estimates:
## mean in group 1 mean in group 2
##      0.7867213      0.6404444
```

There is no evidence that the two groups differ in weight change in the previous six months.

We could still have performed a t-test if the variances were not homogenous by setting the `var.equal` parameter of the `t.test()` function to **FALSE**:

```
t.test(WT ~ LIFE, var.equal = FALSE)
```

```
##
## Welch Two Sample t-test
##
```

```
## data:  WT by LIFE
## t = 0.60608, df = 98.866, p-value = 0.5459
## alternative hypothesis: true difference in means is not equal to 0
## 95 percent confidence interval:
##  -0.3326225  0.6251763
## sample estimates:
## mean in group 1 mean in group 2
##      0.7867213      0.6404444
```

or performed a non-parametric test:

```
wilcox.test(WT ~ LIFE)
```

```
##
##  Wilcoxon rank sum test with continuity correction
##
## data:  WT by LIFE
## W = 1488, p-value = 0.4622
## alternative hypothesis: true location shift is not equal to 0
```

An alternative, and more general, non-parametric test is:

```
kruskal.test(WT ~ LIFE)
```

```
##
##  Kruskal-Wallis rank sum test
##
## data:  WT by LIFE
## Kruskal-Wallis chi-squared = 0.54521, df = 1, p-value = 0.4603
```

We can use the `table()` function to examine the differences in depression between the two groups:

```
table(DEP, LIFE)
```

```
##      LIFE
## DEP  1  2
##    1  0 26
##    2 42 24
##    3 16  1
```

The two distributions look very different from each other. We can test this using a chi-square test on the table:

```
chisq.test(table(DEP, LIFE))
```

```
##
##  Pearson's Chi-squared test
##
## data:  table(DEP, LIFE)
## X-squared = 43.876, df = 2, p-value = 2.968e-10
```

Note that we passed the output of the `table()` function directly to the `chisq.test()` function. We could have saved the table as an object first and then passed the object to the `chisq.test()` function:

```
tab <- table(DEP, LIFE)
chisq.test(tab)
```

```
##
##  Pearson's Chi-squared test
##
```

```
## data:  tab
## X-squared = 43.876, df = 2, p-value = 2.968e-10
```

The `tab` object contains the output of the `table()` function:

```
class(tab)
```

```
## [1] "table"
```

```
tab
```

```
##      LIFE
## DEP   1   2
##    1  0 26
##    2 42 24
##    3 16  1
```

We can pass this table object to another function. For example:

```
fisher.test(tab)
```

```
##
## Fisher's Exact Test for Count Data
##
## data:  tab
## p-value = 1.316e-12
## alternative hypothesis: two.sided
```

When we are finished with the `tab` object we can delete it using the `rm()` function:

```
rm(tab)
```

You can see a list of available objects using the `ls()` function:

```
ls()
```

```
## [1] "fem"
```

This should just show the `fem` object.

We can examine the association between loss of interest in sex and considering suicide in the same way:

```
tab <- table(SEX, LIFE)
tab
```

```
##      LIFE
## SEX   1   2
##    1 58 38
##    2  5 12
```

```
fisher.test(tab)
```

```
##
## Fisher's Exact Test for Count Data
##
## data:  tab
## p-value = 0.03175
## alternative hypothesis: true odds ratio is not equal to 1
## 95 percent confidence interval:
##  1.080298 14.214482
## sample estimates:
## odds ratio
##  3.620646
```

Note that with a two-by-two table the `fisher.test()` function produces an estimate of, and confidence intervals for, the odds ratio. Again, we will delete the `tab` object:

```
rm(tab)
```

We could have performed the Fisher exact test without creating the `tab` object by passing the output of the `table()` function directly to the `fisher.test()` function:

```
fisher.test(table(SEX, LIFE))
```

```
##
## Fisher's Exact Test for Count Data
##
## data:  table(SEX, LIFE)
## p-value = 0.03175
## alternative hypothesis: true odds ratio is not equal to 1
## 95 percent confidence interval:
##  1.080298 14.214482
## sample estimates:
## odds ratio
##  3.620646
```

Choose whichever method you find easiest but remember that it is easy to save the results of any function for later use.


We can explore the correlation between two variables using the `cor()` function:

```
cor(IQ, WT, use = "pairwise.complete.obs")
```

```
## [1] -0.2917158
```

or by using a scatter plot:

```
plot(IQ, WT)
```



and by a formal test:

```
cor.test(IQ, WT)
```

```
##
## Pearson's product-moment correlation
##
## data: IQ and WT
## t = -3.0192, df = 98, p-value = 0.003231
## alternative hypothesis: true correlation is not equal to 0
## 95 percent confidence interval:
## -0.4616804 -0.1010899
## sample estimates:
## cor
## -0.2917158
```

With some functions you can pass an entire data.frame rather than a list of variables:

```
cor(fem, use = "pairwise.complete.obs")
```

```
##          ID          AGE          IQ          ANX          DEP
## ID      1.00000000  0.03069077  0.0370598672 -0.02941825 -0.0554147209
## AGE     0.03069077  1.00000000 -0.4345435680  0.06734300 -0.0387049246
## IQ      0.03705987 -0.43454357  1.0000000000 -0.02323787 -0.0001307404
## ANX     -0.02941825  0.06734300 -0.0232378691  1.00000000  0.5437946347
## DEP     -0.05541472 -0.03870492 -0.0001307404  0.54379463  1.0000000000
## SLP     -0.07268743  0.02606547  0.0812993104  0.22317875  0.5248724551
## SEX      0.08999634  0.10609216 -0.0536558660 -0.21062493 -0.3058422258
## LIFE    -0.05604349 -0.10300193 -0.0915396469 -0.34211268 -0.6139017253
## WT       0.02640131  0.41574411 -0.2917157832  0.11817532  0.0233742465
##          SLP          SEX          LIFE          WT
## ID     -0.072687434  0.08999634 -0.05604349  0.026401310
## AGE     0.026065468  0.10609216 -0.10300193  0.415744109
## IQ      0.081299310 -0.05365587 -0.09153965 -0.291715783
## ANX     0.223178752 -0.21062493 -0.34211268  0.118175321
## DEP     0.524872455 -0.30584223 -0.61390173  0.023374247
## SLP     1.000000000 -0.29053971 -0.35186578 -0.009259774
## SEX    -0.290539709  1.00000000  0.22316967 -0.027826514
## LIFE   -0.351865775  0.22316967  1.00000000 -0.058605326
## WT     -0.009259774 -0.02782651 -0.05860533  1.000000000
```

```
pairs(fem)
```

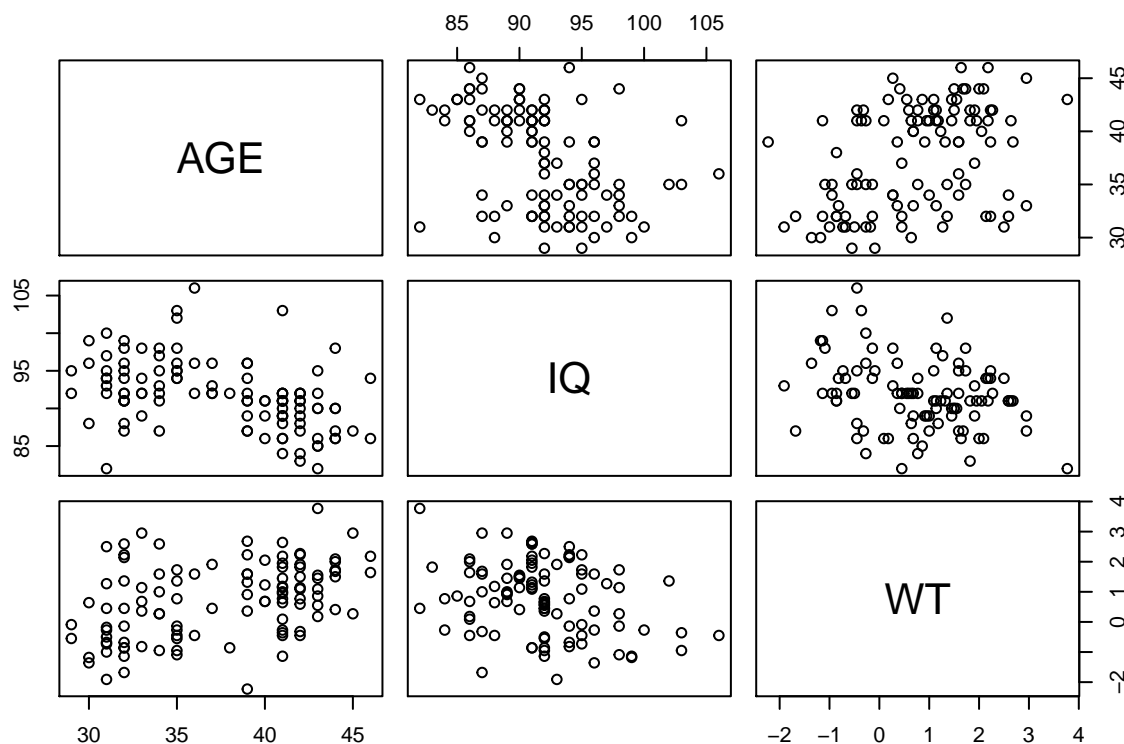


The output can be a little confusing particularly if it includes categorical or record identifying variables. To avoid this we can create a new object that contains only the columns we are interested in using the column binding `cbind()` function:

```
newfem <- cbind(AGE, IQ, WT)
cor(newfem, use = "pairwise.complete.obs")
```

```
##          AGE          IQ          WT
## AGE  1.0000000 -0.4345436  0.4157441
## IQ   -0.4345436  1.0000000 -0.2917158
## WT    0.4157441 -0.2917158  1.0000000
```

```
pairs(newfem)
```



When we have finished with the `newfem` object we can delete it:

```
rm(newfem)
```

There was no real need to create the `newfem` object as we could have fed the output of the `cbind()` function directly to the `cor()` or `pairs()` function:

```
cor(cbind(AGE, IQ, WT), use = "pairwise.complete.obs")
```

```
##           AGE           IQ           WT
## AGE  1.0000000 -0.4345436  0.4157441
## IQ  -0.4345436  1.0000000 -0.2917158
## WT   0.4157441 -0.2917158  1.0000000
```

```
pairs(cbind(AGE, IQ, WT))
```



It is, however, easier to work with the `newfem` object rather than having to retype the `cbind()` function. This is particularly true if you wanted to continue with an analysis of just the three variables.

The relationship between `AGE` and `WT` can be plotted using the `plot()` function:

```
plot(AGE, WT)
```



And tested using the `cor()` and `cor.test()` functions:


```
cor(AGE, WT, use = "pairwise.complete.obs")
```

```
## [1] 0.4157441
```

```
cor.test(AGE, WT)
```

```
##
## Pearson's product-moment correlation
##
## data: AGE and WT
## t = 4.6841, df = 105, p-value = 8.457e-06
## alternative hypothesis: true correlation is not equal to 0
## 95 percent confidence interval:
## 0.2452434 0.5612979
## sample estimates:
## cor
## 0.4157441
```

Or by using the linear modelling `lm()` function:

```
summary(lm(WT ~ AGE))
```

```
##
## Call:
## lm(formula = WT ~ AGE)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -3.10678 -0.85922 -0.05453  0.71434  2.70874
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -3.25405     0.85547  -3.804  0.00024 ***
## AGE          0.10592     0.02261   4.684 8.46e-06 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 1.128 on 105 degrees of freedom
## (11 observations deleted due to missingness)
## Multiple R-squared:  0.1728, Adjusted R-squared:  0.165
## F-statistic: 21.94 on 1 and 105 DF, p-value: 8.457e-06
```

We use the `summary()` function here to extract summary information from the output of the `lm()` function.

It is often more useful to use `lm()` to create an object:

```
fem.lm <- lm(WT ~ AGE)
```


And use the output in other functions:

```
summary(fem.lm)
```

```
##
## Call:
## lm(formula = WT ~ AGE)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
```

```
## -3.10678 -0.85922 -0.05453  0.71434  2.70874
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)
## (Intercept) -3.25405    0.85547  -3.804  0.00024 ***
## AGE          0.10592    0.02261   4.684  8.46e-06 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 1.128 on 105 degrees of freedom
## (11 observations deleted due to missingness)
## Multiple R-squared:  0.1728, Adjusted R-squared:  0.165
## F-statistic: 21.94 on 1 and 105 DF,  p-value: 8.457e-06



plot(AGE, WT)
abline(fem.lm)
```




In this case we are passing the intercept and slope information held in the `fem.lm` object to the `abline()` function which draws a regression line. The `abline()` function adds to an existing plot. This means that you need to keep the scatter plot of AGE and WT open before issuing the `abline()` function call.

A useful function to apply to the `fem.lm` object is `plot()` which produces diagnostic plots of the linear model:

```
plot(fem.lm)
```






Objects created by the `lm()` function (or any of the modelling functions) can use up a lot of memory so we should remove them when we no longer need them:

```
rm(fem.lm)
```

It might be interesting to see whether a similar relationship exists between **AGE** and **WT** for those who have and have not considered suicide. This can be done using the `coplot()` function:

```
coplot(WT ~ AGE | as.factor(LIFE))
```

Given : as.factor(LIFE)




```
##
```

```
## Missing rows: 21, 22, 31, 43, 44, 45, 69, 81, 101, 104, 114, 115
```

The two plots look similar. We could also use `coplot()` to investigate the relationship between `AGE` and `WT` for categories of both `LIFE` and `SEX`:

```
coplot(WT ~ AGE | as.factor(LIFE) * as.factor(SEX))
```



```
##
## Missing rows: 12, 17, 21, 22, 31, 43, 44, 45, 66, 69, 81, 101, 104, 105, 114, 115
```

although the numbers are too small for this to be useful here.

We used the `as.factor()` function with the `coplot()` function to ensure that R was aware that the `LIFE` and `SEX` columns hold categorical data.

We can check the way variables are stored using the `data.class()` function:

```
data.class(fem$SEX)
```

```
## [1] "numeric"
```

We can ‘apply’ this function to all columns in a data.frame using the `sapply()` function:

```
sapply(fem, data.class)
```

```
##      ID      AGE      IQ      ANX      DEP      SLP      SEX
## "numeric" "numeric" "numeric" "numeric" "numeric" "numeric" "numeric"
##      LIFE      WT
## "numeric" "numeric"
```

The `sapply()` function is part of a group of functions that apply a specified function to data objects:

Function(s)	Applies a function to ...
<code>apply()</code>	rows and columns of matrices, arrays, and tables
<code>lapply()</code>	components of lists and data.frames
<code>sapply()</code>	components of lists and data.frames
<code>mapply()</code>	components of lists and data.frames
<code>tapply()</code>	subsets of data

Related functions are `aggregate()` which compute summary statistics for subsets of data, `by()` which applies a function to a data.frame split by factors, and `sweep()` which applies a function to an array.

The parameters of most R functions have default values. These are usually the most used and most useful parameter values for each function. The `cor.test()` function, for example, calculates *Pearson's product moment correlation coefficient* by default. This is an appropriate measure for data from a bivariate normal distribution. The DEP and ANX variables contain ordered data. An appropriate measure of correlation between DEP and ANX is *Kendall's tau*. This can be obtained using:

```
cor.test(DEP, ANX, method = "kendall")

##
## Kendall's rank correlation tau
##
## data: DEP and ANX
## z = 5.5606, p-value = 2.689e-08
## alternative hypothesis: true tau is not equal to 0
## sample estimates:
##      tau
## 0.4950723
```

Before we finish we should save the `fem` data.frame so that next time we want to use it we will not have to bother with recoding the missing values to the special NA value. This is done with the `write.table()` function:

```
write.table(fem, file = "newfem.dat", row.names = FALSE)
```

Everything in R is either a function or an object. Even the command to quit R is a function:

```
q()
```

When you call the `q()` function you will be asked if you want to save the workspace image. If you save the workspace image then all of the objects and functions currently available to you will be saved. These will then be automatically restored the next time you start R in the current working directory.

For this exercise there is no need to save the workspace image so click the **No or Don't Save** button (GUI) or enter `n` when prompted to save the workspace image (terminal).

2.1 Summary

- R is a functional system. Everything is done by calling functions.
- R provides a large set of functions for descriptive statistics, charting, and statistical inference.
- Functions can be chained together so that the output of one function is the input of another function.
- R is an object oriented system. We can use functions to create objects that can then be manipulated or passed to other functions for subsequent analysis.

Chapter 3

Manipulating objects and creating new functions

In this exercise we will explore how to manipulate R objects and how to write functions that can manipulate and extract data and information from R objects and produce useful analyses.

Before we go any further we should start R and retrieve a dataset:

```
salex <- read.table("salex.dat", header = TRUE, na.strings = "9")
```

Missing values are coded as 9 throughout this dataset so we can use the `na.strings` parameter of the `read.table()` function to replace all 9's with the special NA code when we retrieve the dataset. Check that this works by examining the `salex` data.frame:

```
salex
```

##	ILL	HAM	BEEF	EGGS	MUSHROOM	PEPPER	PORKPIE	PASTA	RICE	LETTUCE	TOMATO
## 1	1	1	1	1	1	1	2	2	2	2	2
## 2	1	1	1	1	2	2	1	2	2	2	1
## 3	1	1	1	1	1	1	1	1	1	1	2
## 4	1	1	1	1	2	2	2	2	2	1	1
## 5	1	1	1	1	1	1	1	1	1	1	1
## 6	1	1	1	1	2	2	2	2	2	2	1
## 7	1	1	1	1	1	1	1	2	2	2	2
## 8	1	1	2	1	1	1	2	1	1	1	2
## 9	1	1	1	1	2	1	1	2	1	2	2
## 10	1	1	1	1	2	1	1	1	1	1	1
## 11	1	2	2	1	1	1	2	2	2	1	1
## 12	1	1	1	1	2	2	2	2	2	2	2
## 13	2	2	1	2	2	2	1	2	2	2	1
## 14	1	1	1	1	2	2	2	1	1	2	1
## 15	1	1	1	1	1	1	2	1	1	2	2
## 16	1	1	1	1	1	1	1	2	2	2	2
## 17	1	1	1	1	1	1	1	1	1	1	1
## 18	2	1	1	2	2	2	2	2	2	2	2
## 19	2	1	1	1	1	2	2	1	1	2	1
## 20	2	1	1	2	2	2	2	2	2	2	2
## 21	2	2	2	2	2	2	2	2	2	2	2
## 22	1	1	1	1	2	2	2	2	2	1	1
## 23	1	2	1	2	2	2	2	1	1	2	1

## 24	1	1	1	1	2	1	2	1	1	2	2
## 25	1	1	1	2	1	1	1	1	1	1	1
## 26	1	1	2	1	1	1	2	2	2	1	1
## 27	1	1	1	1	2	2	1	2	1	1	1
## 28	1	1	1	1	1	1	2	1	1	2	2
## 29	1	2	1	1	1	NA	2	1	1	1	1
## 30	1	1	1	2	2	2	1	2	2	2	2
## 31	1	1	1	1	1	2	2	1	1	2	2
## 32	1	1	1	1	1	2	NA	2	1	1	1
## 33	1	1	1	1	2	2	2	1	2	2	2
## 34	1	1	1	1	1	2	2	2	2	1	1
## 35	1	1	1	1	1	1	1	1	2	2	1
## 36	2	2	1	2	2	2	2	2	2	2	2
## 37	1	1	1	1	1	1	2	1	1	1	1
## 38	1	1	1	2	2	2	1	1	1	1	2
## 39	1	1	1	1	1	1	1	2	2	1	2
## 40	1	1	1	1	1	1	1	2	2	1	1
## 41	1	1	1	2	2	1	2	1	1	1	1
## 42	1	1	1	2	2	2	2	2	2	2	2
## 43	1	1	1	1	1	1	2	1	1	1	1
## 44	1	2	1	2	2	2	1	2	2	1	2
## 45	1	1	1	1	1	2	2	2	1	1	1
## 46	1	1	1	2	2	2	2	1	1	1	1
## 47	1	1	1	1	2	2	2	2	1	1	2
## 48	1	1	1	1	1	NA	1	1	1	2	2
## 49	1	1	1	1	2	1	2	2	1	1	1
## 50	1	2	1	1	2	2	2	1	2	2	1
## 51	2	2	1	2	2	2	2	2	2	2	2
## 52	2	1	1	2	2	2	2	1	2	2	1
## 53	2	1	1	2	2	2	1	2	2	2	1
## 54	2	1	1	2	1	2	1	2	2	2	1
## 55	2	1	1	1	1	1	2	2	1	2	2
## 56	2	1	1	2	2	2	2	2	2	2	1
## 57	2	1	1	1	1	1	1	2	2	2	2
## 58	2	1	1	1	2	2	1	2	1	2	2
## 59	2	1	1	2	2	2	2	2	2	2	2
## 60	2	2	2	2	2	2	1	2	2	2	2
## 61	2	1	1	2	2	2	1	2	2	2	2
## 62	2	1	2	2	2	2	2	2	2	1	1
## 63	1	1	1	1	1	1	2	2	2	2	1
## 64	2	1	1	2	2	2	2	2	2	2	2
## 65	2	1	1	1	1	2	1	2	1	2	2
## 66	2	2	1	2	2	2	2	2	2	2	2
## 67	2	2	1	2	2	2	2	2	2	2	2
## 68	2	1	1	2	1	1	1	1	2	2	1
## 69	2	2	1	2	2	2	2	2	2	2	2
## 70	2	2	1	2	2	2	2	2	2	2	2
## 71	1	1	2	2	2	2	1	2	1	2	2
## 72	2	1	2	1	NA	NA	2	2	2	2	1
## 73	1	1	1	1	2	2	1	2	2	2	2
## 74	1	1	2	1	NA	NA	2	1	1	1	1
## 75	1	1	2	2	2	1	2	1	2	1	1
## 76	1	1	1	1	2	2	1	1	2	2	2
## 77	1	1	1	NA	NA	NA	1	2	1	1	1

##	COLESLAW	CRISPS	PEACHCAKE	CHOCOLATE	FRUIT	TRIFLE	ALMONDS
## 1	2	2	2	2	2	2	2
## 2	2	2	2	2	2	2	2
## 3	2	1	2	1	2	2	2
## 4	2	2	2	1	2	2	2
## 5	1	2	2	1	2	1	2
## 6	1	1	2	1	2	2	2
## 7	1	1	1	2	2	2	2
## 8	1	1	2	2	2	1	2
## 9	2	2	2	2	2	1	2
## 10	1	1	2	2	2	1	1
## 11	2	2	2	2	2	2	NA
## 12	2	1	2	1	2	2	2
## 13	2	1	2	2	1	2	NA
## 14	1	1	2	2	2	1	2
## 15	1	1	2	2	2	1	1
## 16	1	2	2	2	2	2	2
## 17	1	2	2	2	2	2	2
## 18	2	2	2	2	2	2	2
## 19	1	1	2	2	1	2	2
## 20	2	2	2	1	2	2	2
## 21	2	2	2	2	2	2	2
## 22	2	1	2	1	2	2	2
## 23	1	2	2	2	2	2	NA
## 24	1	1	2	2	2	1	2
## 25	1	2	2	2	2	1	NA
## 26	1	2	2	2	2	1	2
## 27	1	1	1	1	2	1	2
## 28	2	1	2	2	2	2	NA
## 29	1	1	2	2	2	2	NA
## 30	2	2	2	2	2	2	2
## 31	2	2	2	2	2	2	2
## 32	2	2	2	2	2	2	2
## 33	1	2	2	2	2	2	2
## 34	1	2	2	2	2	1	2
## 35	1	2	2	2	2	1	2
## 36	2	2	2	2	2	2	NA
## 37	1	1	2	1	2	1	2
## 38	2	2	2	2	2	2	2
## 39	2	2	2	1	2	2	2
## 40	1	2	2	2	2	2	2
## 41	1	1	2	2	NA	1	NA
## 42	2	2	2	2	2	2	NA
## 43	1	1	2	2	2	2	NA
## 44	2	2	2	2	2	2	2
## 45	1	2	2	2	2	1	2
## 46	1	2	2	2	2	1	2
## 47	2	2	2	NA	2	1	2
## 48	2	1	2	2	2	2	2
## 49	1	1	2	2	2	1	2
## 50	NA	2	2	1	2	1	1
## 51	2	2	2	2	2	2	NA
## 52	2	2	2	1	2	2	1
## 53	2	2	2	2	1	2	2

```
## 54      2      2      2      2      2      2      2
## 55      2      1      2      2      2      2      2
## 56      2      2      2      2      1      2      2
## 57      1      2      2      2      2      2      1
## 58      2      1      1      2      2      2      2
## 59      2      1      1      2      2      1      2
## 60      2      2      2      2      2      2      2
## 61      2      1      2      2      2      1      1
## 62      2      1      2      2      2      2      2
## 63      1      2      2      1      1      2      2
## 64      2      2      2      2      2      2      2
## 65      1      1      2      2      2      1      2
## 66      2      2      2      2      2      1      NA
## 67      2      1      2      2      2      2      2
## 68      2      2      2      2      2      2      2
## 69      2      1      2      2      2      2      2
## 70      2      2      2      2      2      2      2
## 71      2      2      2      2      2      2      2
## 72      2      2      2      2      2      1      2
## 73      2      2      2      2      2      2      2
## 74      1      1      2      1      2      2      2
## 75      1      1      2      2      2      2      NA
## 76      2      2      2      2      2      2      NA
## 77      1      1      2      2      2      2      2
```

```
names(salex)
```

```
## [1] "ILL"      "HAM"      "BEEF"      "EGGS"      "MUSHROOM"
## [6] "PEPPER"   "PORKPIE"   "PASTA"      "RICE"      "LETTUCE"
## [11] "TOMATO"   "COLESLAW" "CRISPS"     "PEACHCAKE" "CHOCOLATE"
## [16] "FRUIT"    "TRIFLE"    "ALMONDS"
```

This data comes from a food-borne outbreak. On Saturday 17th October 1992, eighty-two people attended a buffet meal at a sports club. Within fourteen to twenty-four hours, fifty-one of the participants developed diarrhoea, with nausea, vomiting, abdominal pain and fever.

The columns in the dataset are as follows:

ILL	Ill or not-ill
HAM	Baked ham
BEEF	Roast beef
EGGS	Eggs
MUSHROOM	Mushroom flan
PEPPER	Pepper flan
PORKPIE	Pork pie
PASTA	Pasta salad
RICE	Rice salad
LETTUCE	Lettuce
TOMATO	Tomato salad
COLESLAW	Coleslaw
CRISPS	Crisps
PEACHCAKE	Peach cake
CHOCOLATE	Chocolate cake
FRUIT	Tropical fruit salad
TRIFLE	Trifle
ALMONDS	Almonds

Data is available for seventy-seven of the eighty-two people who attended the sports club buffet. All of the variables are coded 1=yes, 2=no.

We can use the `attach()` function to make it easier to access our data:

```
attach(salex)
```

```
## The following objects are masked from salex (pos = 4):
##
##      ALMONDS, BEEF, CHOCOLATE, COLESLAW, CRISPS, EGGS, FRUIT, HAM,
##      ILL, LETTUCE, MUSHROOM, PASTA, PEACHCAKE, PEPPER, PORKPIE,
##      RICE, TOMATO, TRIFLE
```

The two-by-two table is a basic epidemiological tool. In analysing data from a food-borne outbreak collected as a retrospective cohort study, for example, we would tabulate each exposure (suspect foodstuffs) against the outcome (illness) and calculate risk ratios and confidence intervals. R has no explicit function to calculate risk ratios from two-by-two tables but we can easily write one ourselves.

The first step in writing such a function would be to create the two-by-two table. This can be done with the `table()` function. We will use a table of HAM by ILL as an illustration:

```
table(HAM, ILL)
```

This command produces the following output:

```
##      ILL
## HAM   1   2
##    1 46 17
##    2   5   9
```

We can manipulate the output directly but it is easier if we instruct R to save the output of the `table()` function in an object:

```
tab <- table(HAM, ILL)
```

The `tab` object contains the output of the `table()` function:

```
tab
```

```
##      ILL
## HAM   1   2
##    1 46 17
##    2   5   9
```

As it is stored in an object we can examine its contents on an item by item basis.

The `tab` object is an object of class `table`:

```
class(tab)
```

```
## [1] "table"
```

We can extract data from a table object by using indices or row and column co-ordinates:

```
tab[1,1]
```

```
## [1] 46
```

```
tab[1,2]
```

```
## [1] 17
```

```
tab[2,1]
```

```
## [1] 5
```

The numbers in the square brackets refer to the *position* (as row and column co-ordinates) of the data item in the table *not* the *values* of the variables. We can extract data using the values of the row and column variables by enclosing the index values in double quotes (“). For example:

```
tab["1", "1"]
```

```
## [1] 46
```

The two methods of extracting data may be combined. For example:

```
tab[1, "1"]
```

```
## [1] 46
```

We can calculate a risk ratio using the extracted data:

```
(tab[1,1]/(tab[1,1]+tab[1,2]))/(tab[2,1]/(tab[2,1]+tab[2,2]))
```

Which returns a risk ratio of

```
## [1] 2.044444
```

This is a tedious calculation to have to type in every time you need to calculate a risk ratio from a two-by-two table. It would be better to have a function that calculates and displays the risk ratio automatically. Fortunately, R allows us to do just that.